HW2

Speaker：Alan

Advisor：Prof. An-Yeu Wu

Date：2022/09/20

Roofline Visual Performance Model

- A simple performance model : Execution vs. Data Transfer
- Three key factors
  - System Spec (Hardware Level)
    - Computation : peak floating point performance
      - Floating-point ops /sec
    - Memory: peak memory bandwidth
      - Bytes per sec
  - Program characteristics (Algorithm Level)
    - Arithmetic intensity : Floating-Point Ops/ byte
      - Ratio of floating-point operations in a program to the number of data types accessed by a program from main memory

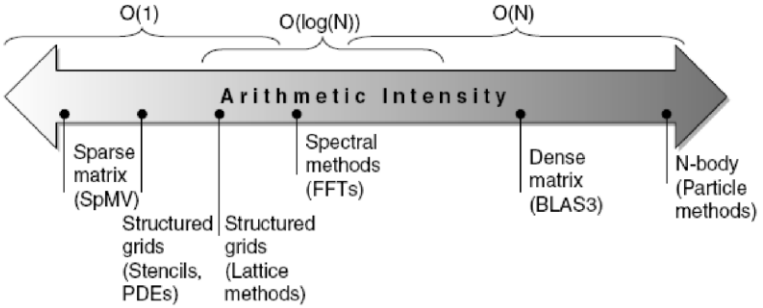

Simplistic view

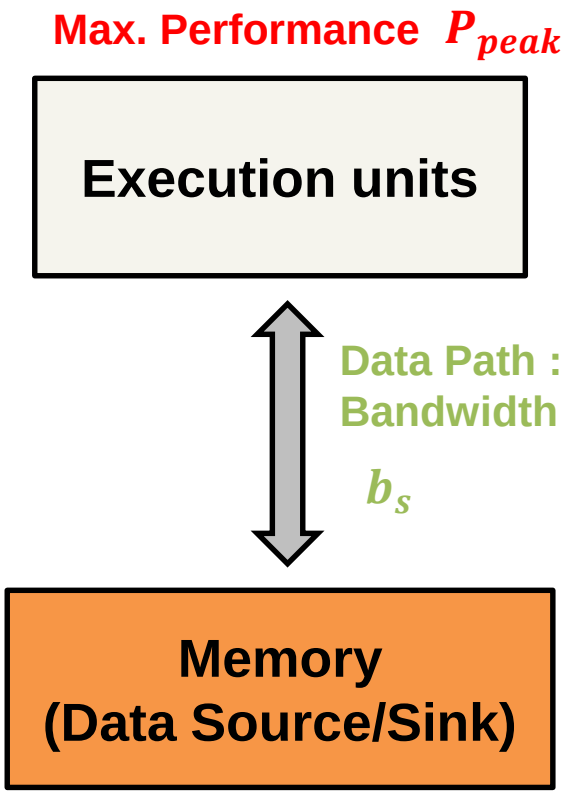

- How fast can tasks be processed?
- The Bottleneck:
  - The execution of Work:
    - P_{peak} [flops : flop/s]
  - The data path
    - $I * b_s$ [flop/byte * byte/s]
    - I : Arithmetic Intensity
- Roofline Model Equation

$$P = \min(P_{peak}, I * bs)$$

Roofline Diagram

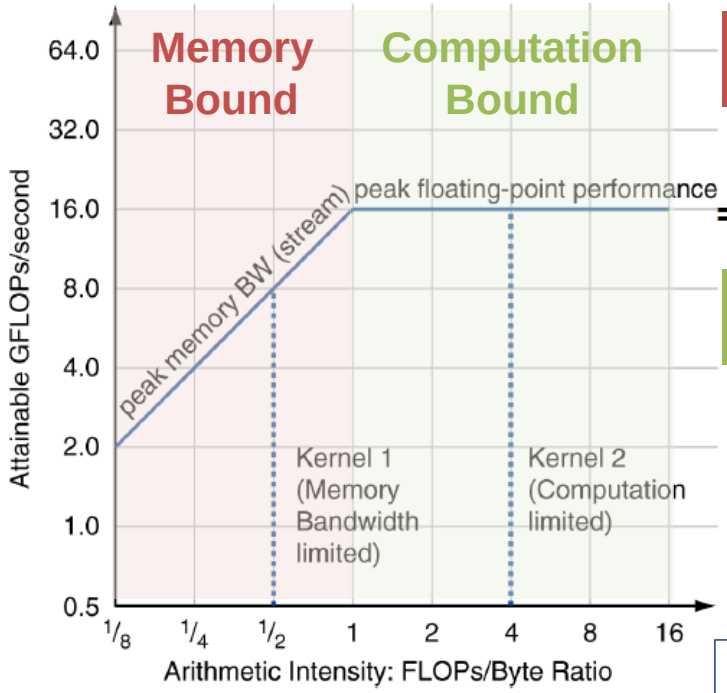

Memory Bandwidth Limited

Floating-Point Ops/sec

= Peak Bytes/sec x Floating-Point Ops/ byte

Computation Limited

Peak Floating-Point Ops/sec

(assuming 16GB/sec peak bandwidth)

Roofline Analysis

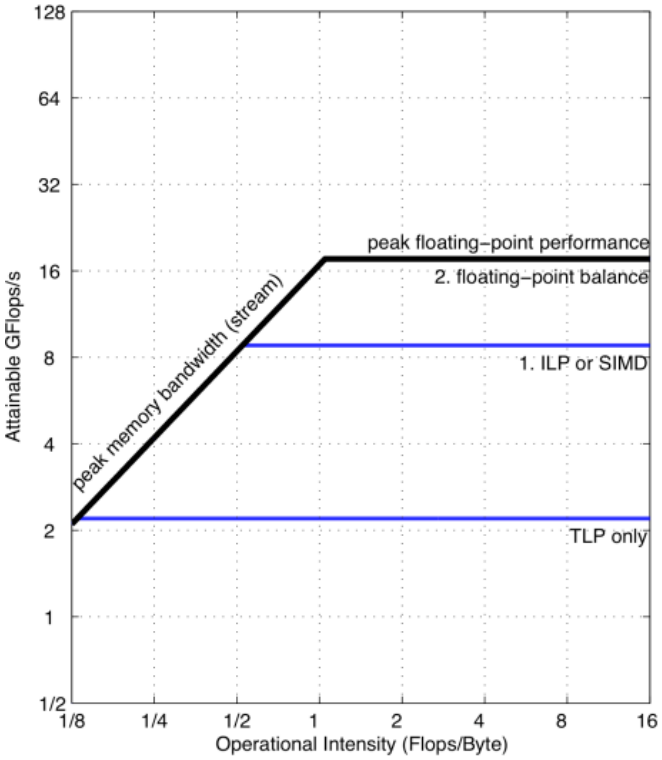

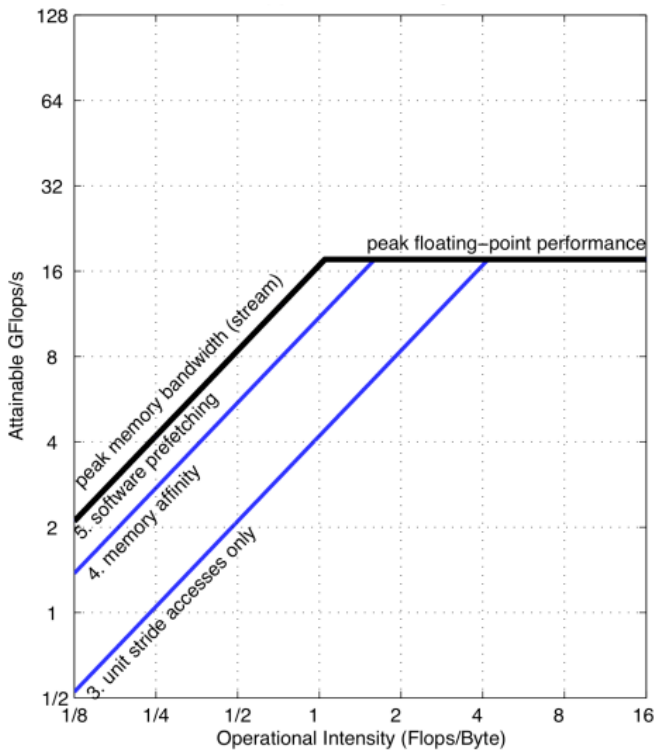

- The Roofline model gives an upper bound to performance.
  - Need some techniques to achieve the ceiling

*TLP : Thread-Level Parallelism
ILP : Instruction-Level Parallelism

HW2 – Roofline Model

- TPU Example

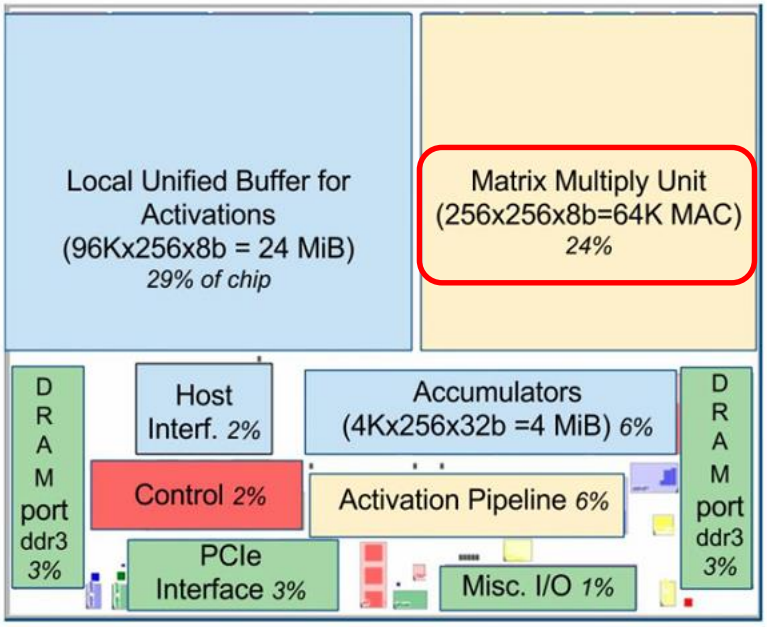

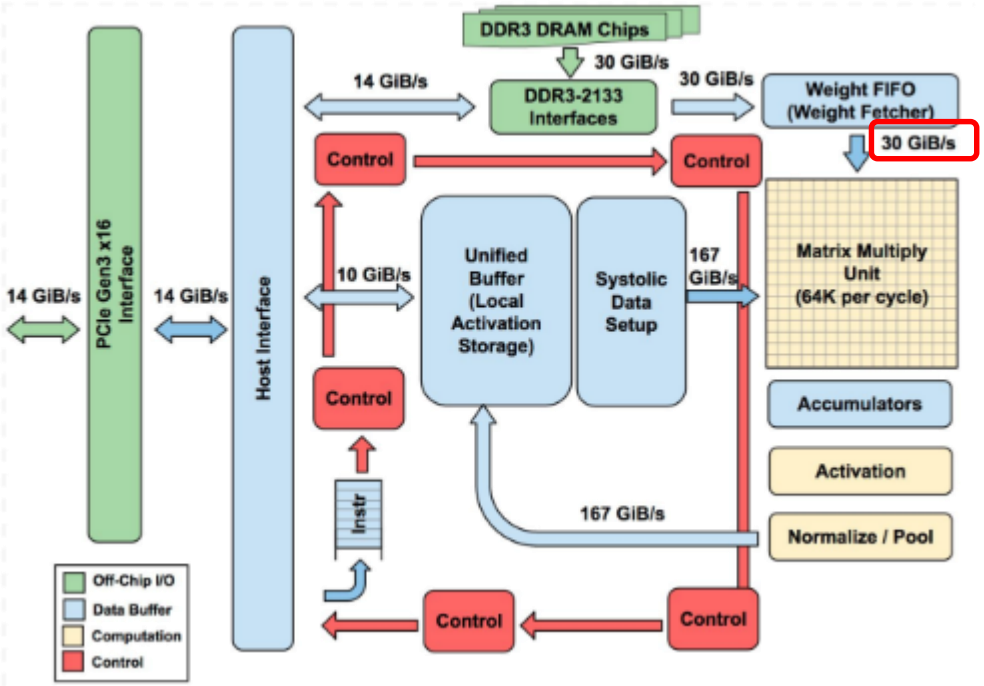

Peak Performance :

$$\underbrace{64 * 1024}_{\text{Number of PE}} * \underbrace{2}_{\text{MAC}} * \underbrace{700 * 10^6}_{\text{Frequency}} = 91.7504\ Tops$$

Bandwidth:

$$30 * \frac{2^{30}}{10^9} = 32.2 GB/s$$

HW2 – Roofline Model

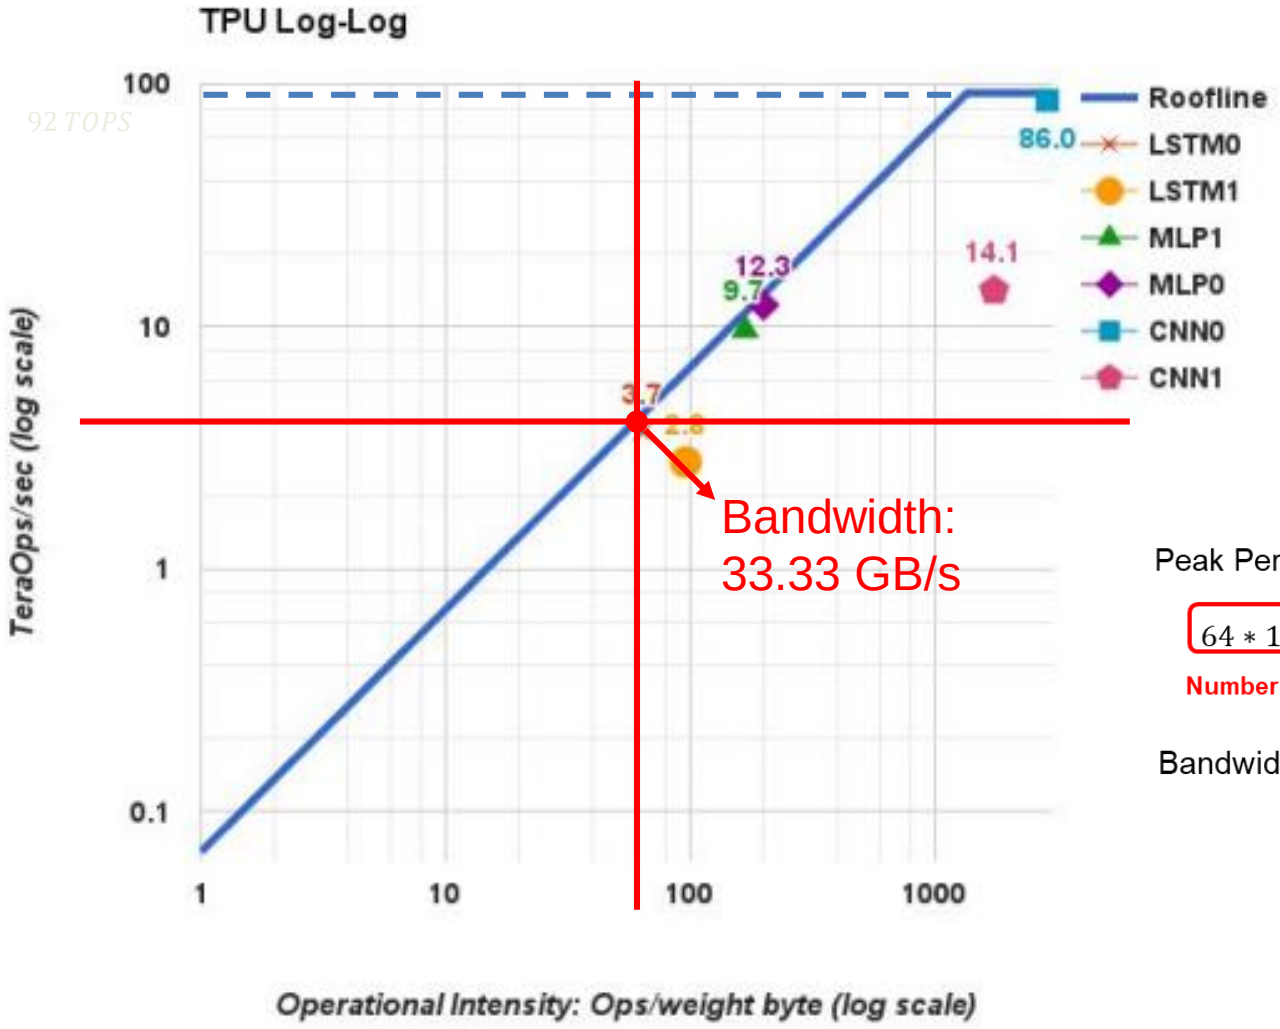

Peak Performance :

$$\underbrace{64 * 1024}_{\text{Number of PE}} * \underbrace{2}_{\text{MAC}} * \underbrace{700 * 10^6}_{\text{Frequency}} = 91.7504\ Tops$$

Bandwidth:

$$30 * \frac{2^{30}}{10^9} = 32.2 GB/s$$

Roofline Model : TPU

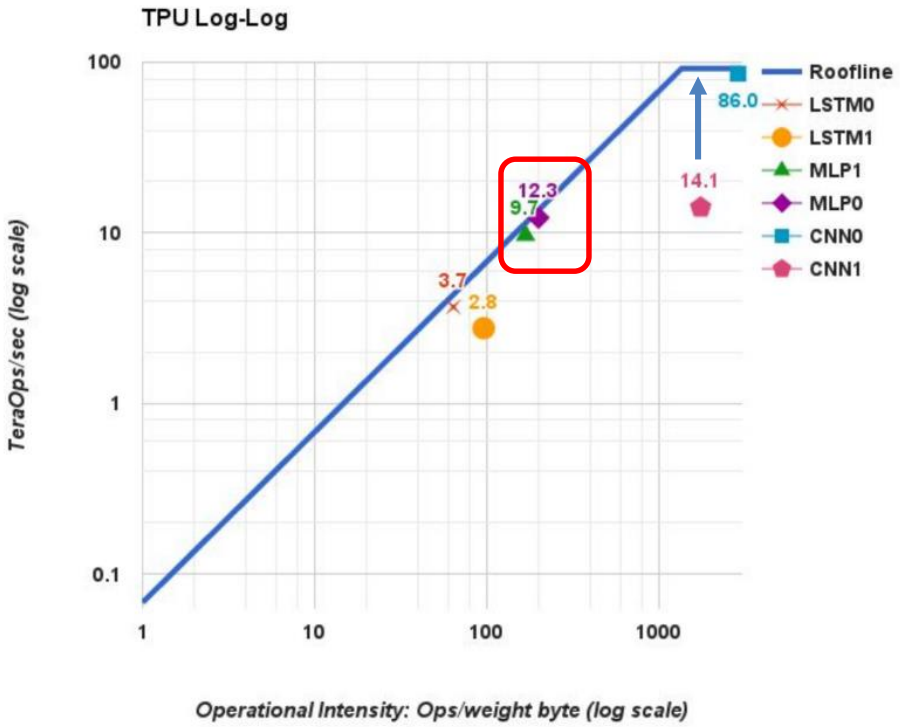

Application	*CNN0*	*CNN1*
Array active cycles	78.2%	46.2%
Useful MACs in 64K matrix (% peak)	78.2%	22.5%
Unused MACs	0.0%	23.7%
Weight stall cycles	0.0%	28.1%
Weight shift cycles	0.0%	7.0%
Non-matrix cycles	21.8%	18.7%
RAW stalls	3.5%	22.8%
Input data stalls	3.4%	0.6%
TeraOps/sec (92 Peak)	86.0	14.1

Low Utilization

- Fully connected layer is less operation-intensive than convolution layer.
- CNN1 has some layers with shallow feature depths.
  - Utilization is not high
    - The actual efficiency is far away from ceiling

Problem 1 (20 points)

- Recap the concept of arithmetic intensity (AI)
  - Compare and discuss the arithmetic intensity of *LSTM*, *MLP*, and *CNN*

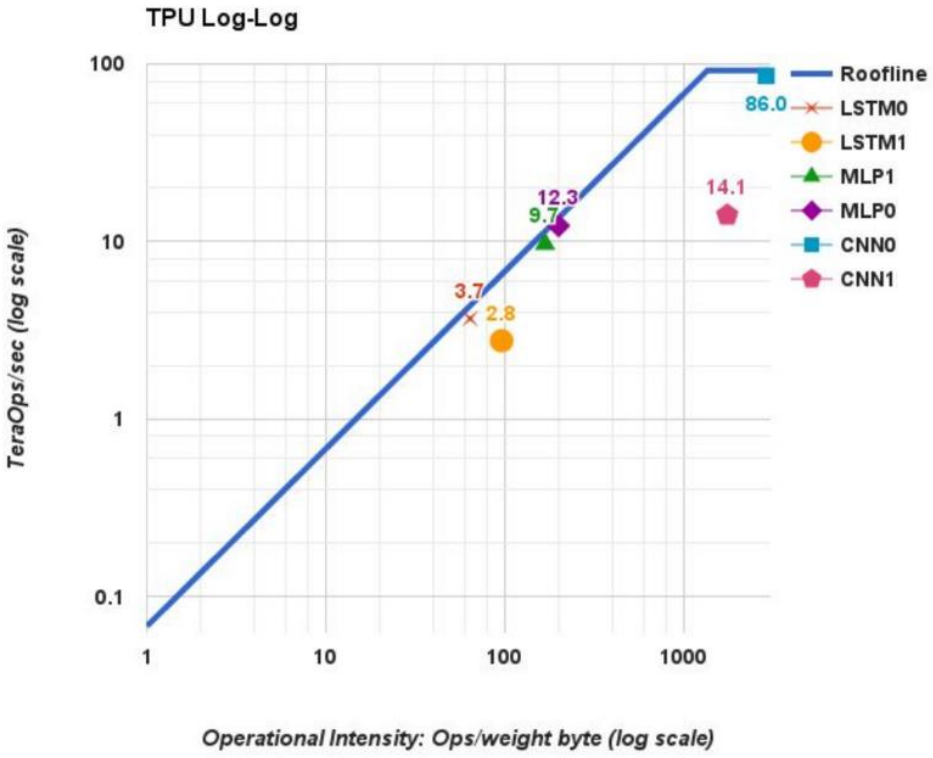

Problem 2 (20 points)

- Goal
  - Understand the meaning of the roofline model
  - Adjust the roofline according to different specification
- Plot the roofline curve if the TPU has upgraded its PE array
  - From 256x256 to 320x320

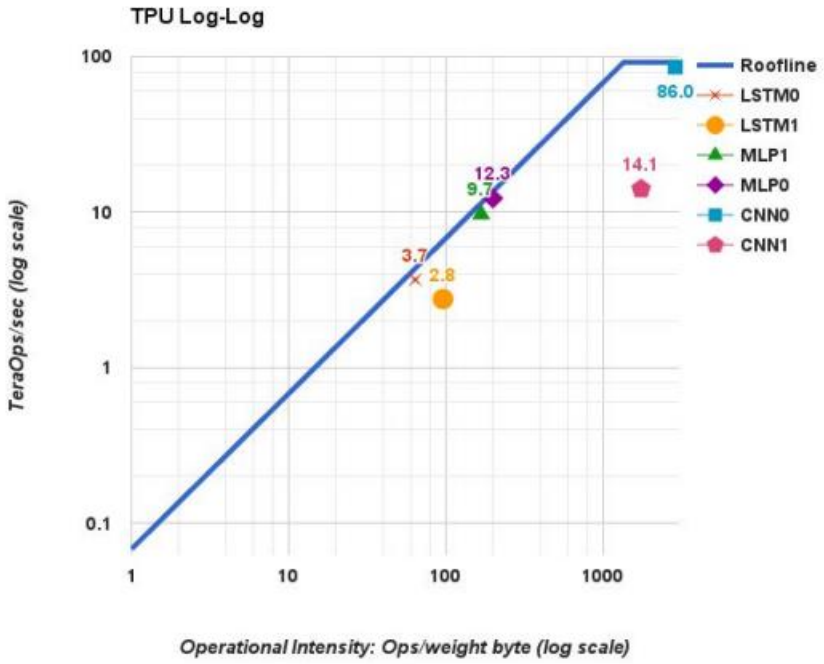

Problem 3 (60 points)

- Goal
  - Understand the meaning of the roofline model
  - Adjust the roofline according to different specification
- Plot the roofline model if we change the hardware
  - CPU (Haswell), GPU (Nvidia K80), and TPU

Model	Die										Benchmarked Servers				
	mm^2	nm	MHz	TDP	Measured		TOPS/s		GB/s	On-Chip Memory	Dies	DRAM Size	TDP	Measured	
					Idle	Busy	8b	FP						Idle	Busy
Haswell E5-2699 v3	662	22	2300	145W	41W	145W	2.6	1.3	51	51 MiB	2	256 GiB	504W	159W	455W
NVIDIA K80 (2 dies/card)	561	28	560	150W	25W	98W	--	2.8	160	8 MiB	8	256 GiB (host) + 12 GiB x 8	1838W	357W	991W
TPU	NA*	28	700	75W	28W	40W	92	--	34	28 MiB	4	256 GiB (host) + 8 GiB x 4	861W	290W	384W

Requirements

- The report should be merged as a single **pdf file** and **uploaded to NTU COOL**.
  - Example of filename: AVLSI_HW2_d09943011.pdf
  - Note that you have to replace d09943011 with your student ID number
- Deadline: 2022/09/26 23:59
  - Late submission will only get half score (deadline: 2022/09/30 23:59)